



**KOÇ  
UNIVERSITY**

# CPU Scheduling

Hakan Ayrıl

Lecture 6

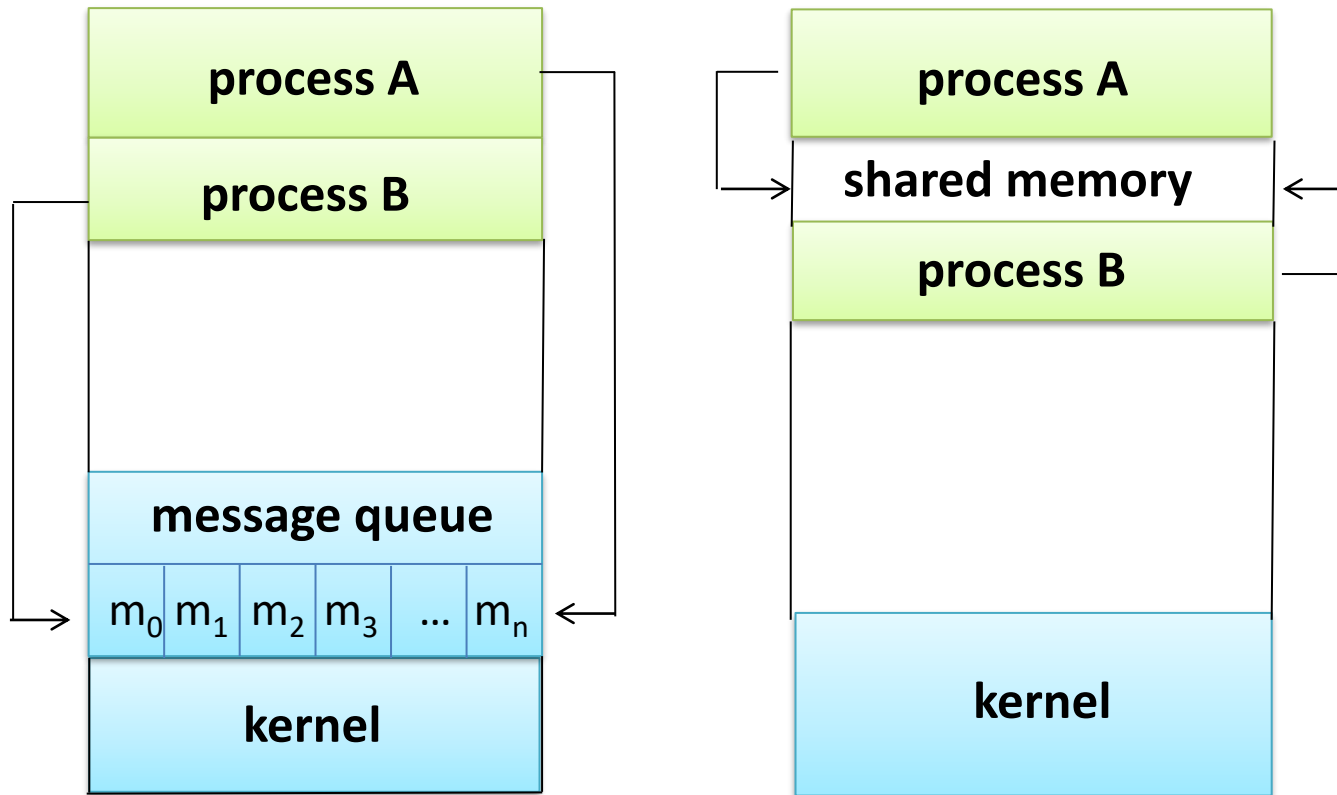
COMP304 - Operating Systems (OS)

# Inter-process Communication (IPC)

- *An independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* processes can affect or be affected by the execution of another processes
- Cooperating processes need **inter-process communication**
- Two models of IPC
  - **Shared memory**
  - **Message passing**

# Two Models of Communication

Message Passing vs Shared Memory



- Message passing requires the message of A to be copied to a buffer and copied to process B's memory – thus it is slower but safer

# Scheduling

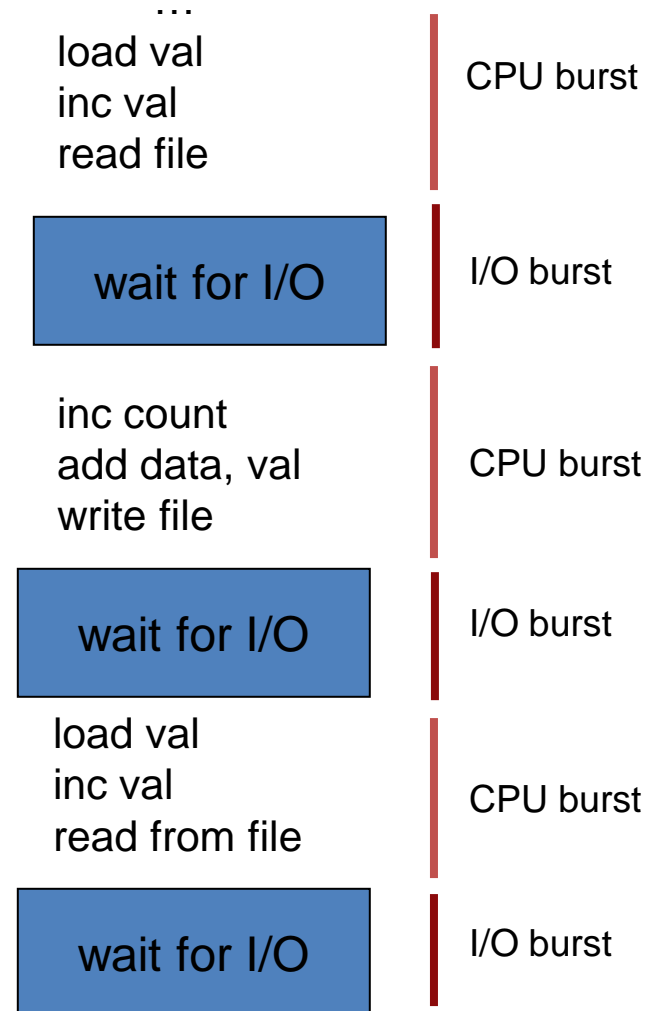
- One of the main tasks of an OS is to schedule processes to execute.
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue.
- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU to that process.

# Schedulers

- **Short-term scheduler** is invoked very frequently (milliseconds)
  - $\Rightarrow$  must be fast.
- **Long-term scheduler** is invoked very infrequently (seconds, minutes)
  - $\Rightarrow$  can be slow.
- The long-term scheduler controls the ***degree of multiprogramming***.

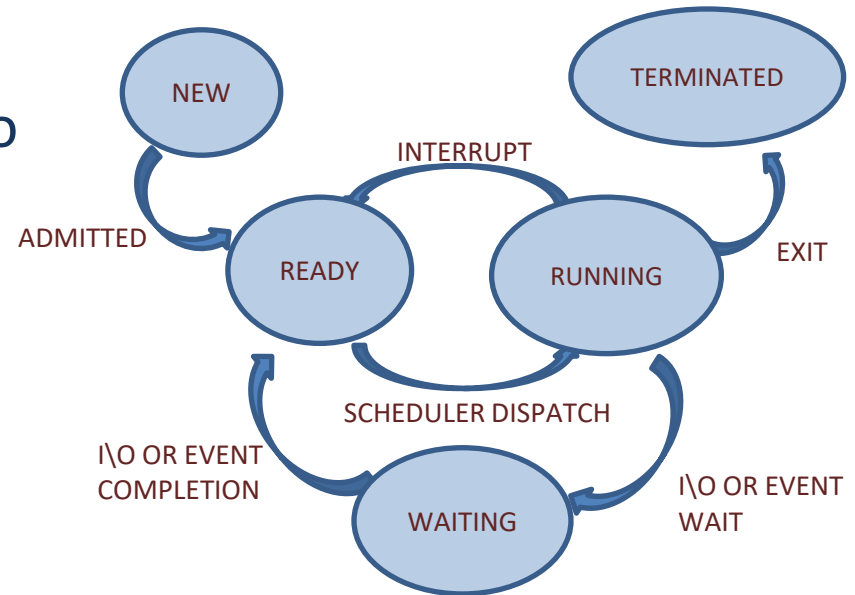
# CPU-I/O Burst Cycle

- **CPU-I/O Burst Cycle** – Process execution consists of a *cycle* of CPU execution and I/O wait
- Processes can be described as either:
  - *I/O-bound process* – spends more time doing I/O than computations; many, short CPU bursts.
  - *CPU-bound process* – spends more time doing computations; few, very long CPU bursts.



# CPU Scheduler

- Selects among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state.
  2. Switches from running to ready state.
  3. Switches from waiting to ready.
  4. Terminates.



# (Non)-preemptive

- Non-preemptive
    - Process voluntarily releases CPU
  - Preemptive
    - OS kicks the process out from the CPU
- 
- 1 and 4 non-preemptive
  - 2 and 3 are preemptive





# Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that completes their execution per time unit
- **Turnaround time** – amount of time to execute a particular process (time between entry and exit)
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

# Scheduling

- Let  $P = \{p_i \mid 0 \leq i < n\}$  = set of processes
- Let  $S(p_i) \in \{\text{running, ready, waiting}\}$
- Let  $t(p_i)$  = Time process needs to be in running state (the service time, CPU burst)
- Let  $T_{\text{TRnd}}(p_i)$  = Time from  $p_i$  first enters ready to last exits system (turnaround time)
- Batch Throughput rate = inverse of avg  $T_{\text{TRnd}}$
- Let  $R(p_i)$  = Time  $p_i$  is in ready state before first transition to running (or response time) (different than “waiting time”)

# Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

It is important to minimize the variance in response time than minimize the average response time – provides fairness

# Dispatcher

- **Dispatcher** module is part of the OS that gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start another running.


# Scheduling Algorithms

If you were the OS, how would you decide who should run next?

# First-Come, First Served (FCFS)

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1$  ,  $P_2$  ,  $P_3$   
The Gantt Chart for the schedule is:




- Waiting times for:  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order


$P_2, P_3, P_1$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case.
- Convoy effect:** short process behind a long process

# FCFS Scheduling (Cont.)




$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$R(p_0) = 0$$



# FCFS Scheduling (Cont.)



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$


$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$R(p_0) = 0$$

$$R(p_1) = T_{\text{TRnd}}(p_0) = 350$$

# FCFS Scheduling (Cont.)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$T_{\text{TRnd}}(p_2) = (\tau(p_2) + T_{\text{TRnd}}(p_1)) = 475 + 475 = 950$$


$$R(p_0) = 0$$

$$R(p_1) = T_{\text{TRnd}}(p_0) = 350$$

$$R(p_2) = T_{\text{TRnd}}(p_1) = 475$$

# FCFS Scheduling (Cont.)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$T_{\text{TRnd}}(p_2) = (\tau(p_2) + T_{\text{TRnd}}(p_1)) = 475 + 475 = 950$$

$$T_{\text{TRnd}}(p_3) = (\tau(p_3) + T_{\text{TRnd}}(p_2)) = 250 + 950 = 1200$$

$$R(p_0) = 0$$


$$R(p_1) = T_{\text{TRnd}}(p_0) = 350$$

$$R(p_2) = T_{\text{TRnd}}(p_1) = 475$$

$$R(p_3) = T_{\text{TRnd}}(p_2) = 950$$

# FCFS Scheduling (Cont.)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$T_{\text{TRnd}}(p_2) = (\tau(p_2) + T_{\text{TRnd}}(p_1)) = 475 + 475 = 950$$

$$T_{\text{TRnd}}(p_3) = (\tau(p_3) + T_{\text{TRnd}}(p_2)) = 250 + 950 = 1200$$

$$T_{\text{TRnd}}(p_4) = (\tau(p_4) + T_{\text{TRnd}}(p_3)) = 75 + 1200 = 1275$$

$$R(p_0) = 0$$

$$R(p_1) = T_{\text{TRnd}}(p_0) = 350$$

$$R(p_2) = T_{\text{TRnd}}(p_1) = 475$$


$$R(p_3) = T_{\text{TRnd}}(p_2) = 950$$

$$R(p_4) = T_{\text{TRnd}}(p_3) = 1200$$

# FCFS Scheduling- Average Wait Time

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

- Easy to implement
- Not a great performer
- Non-preemptive



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$R(p_0) = 0$$

$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$R(p_1) = T_{\text{TRnd}}(p_0) = 350$$

$$T_{\text{TRnd}}(p_2) = (\tau(p_2) + T_{\text{TRnd}}(p_1)) = 475 + 475 = 950$$

$$R(p_2) = T_{\text{TRnd}}(p_1) = 475$$

$$T_{\text{TRnd}}(p_3) = (\tau(p_3) + T_{\text{TRnd}}(p_2)) = 250 + 950 = 1200$$

$$R(p_3) = T_{\text{TRnd}}(p_2) = 950$$

$$T_{\text{TRnd}}(p_4) = (\tau(p_4) + T_{\text{TRnd}}(p_3)) = 75 + 1200 = 1275$$

$$R(p_4) = T_{\text{TRnd}}(p_3) = 1200$$

Average response (wait) time  $R_{\text{avg}} = (0 + 350 + 475 + 950 + 1200) / 5 = 2974 / 5 = 595$


## 2. Shortest Job First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
  - **Non-preemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst.
  - **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This is known as the **Shortest-Remaining-Time-First (SRTF)**.
- SJF is **optimal** – gives minimum average waiting time for a given set of processes.

# Non-preemptive (SJF) Scheduling

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P <sub>1</sub>	0.0	7
P <sub>2</sub>	2.0	4
P <sub>3</sub>	4.0	1
P <sub>4</sub>	5.0	4

- SJF (non-preemptive)




- Average waiting time ?  
 $= (0 + 6 + 3 + 7)/4 = 4$

# Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P <sub>1</sub>	0.0	7
P <sub>2</sub>	2.0	4
P <sub>3</sub>	4.0	1
P <sub>4</sub>	5.0	4

- SJF (preemptive)




- Average waiting time ?  
 $= (9 + 1 + 0 + 2)/4 = 3$



# Example of SJF


$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

# Example of SJF




$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$R(p_1) = 75$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

# Example of SJF



$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$R(p_1) = 75$$


$$T_{\text{TRnd}}(p_3) = \tau(p_3) + \tau(p_1) + \tau(p_4) = 250 + 125 + 75 = 450$$

$$R(p_3) = 200$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

# Example of SJF



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 350 + 250 + 125 + 75 = 800$$

$$R(p_0) = 450$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$R(p_1) = 75$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) + \tau(p_1) + \tau(p_4) = 250 + 125 + 75 = 450$$


$$R(p_3) = 200$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

# Example of SJF

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 350 + 250 + 125 + 75 = 800$$

$$R(p_0) = 450$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$R(p_1) = 75$$

$$T_{\text{TRnd}}(p_2) = \tau(p_2) + \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 475 + 350 + 250 + 125 + 75 = 1275$$

$$R(p_2) = 800$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) + \tau(p_1) + \tau(p_4) = 250 + 125 + 75 = 450$$

$$R(p_3) = 200$$


$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

# Example of SJF

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

- Minimizes waiting time – Why?
- May starve large jobs



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 350 + 250 + 125 + 75 = 800$$

$$R(p_0) = 450$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$R(p_1) = 75$$

$$T_{\text{TRnd}}(p_2) = \tau(p_2) + \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 475 + 350 + 250 + 125 + 75 = 1275$$

$$R(p_2) = 800$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) + \tau(p_1) + \tau(p_4) = 250 + 125 + 75 = 450$$

$$R(p_3) = 200$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$

$$R(p_4) = 0$$

$$\text{Average response (wait) time } R_{\text{avg}} = (450 + 75 + 800 + 200 + 0) / 5 = 1525 / 5 = 305$$

# Shortest Job First

- The SJF is provably optimal
- It gives the minimum average waiting time for a given set of processes
- Moving a short process before a long one decreases the waiting time of the short process more than it increases the waiting time of the long process
- What is the difficulty of using the shortest job first scheduler?

# Determining the Length of the Next CPU Burst

- Can only **estimate the length**.
- Can be done by using the length of previous CPU bursts, using **exponential averaging**.

1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$



# Examples of Exponential Averaging

- $\alpha = 0$

- $\tau_{n+1} = \tau_n$
- Recent information does not count.

- $\alpha = 1$

- $\tau_{n+1} = t_n$
- Only the actual last CPU burst counts.

- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- Since both  $\alpha$  and  $(1 - \alpha)$  are less than or equal to 1, each successive term has less weight than its predecessor.

# Question

\_\_\_\_\_ is the number of processes that are completed per time unit.

- A) CPU utilization
- B) Response time
- C) Turnaround time
- D) Throughput

# Question

- The strategy of making processes that are logically runnable to be temporarily suspended is called :
  - a) Non preemptive scheduling
  - b) Preemptive scheduling
  - c) Shortest job first
  - d) First come First served

Answer is b)

# Reading

- Read Chapter 5
- Read Chapter 4 (Linux Kernel Development)
- Acknowledgments
  - Original slides are by **Didem Unat** which were adapted from
    - Öznur Özkasap (Koç University)
    - Operating System and Concepts (9<sup>th</sup> edition) Wiley

# Puzzle

- A group of people wants to get through a tunnel.
  - A can make it in 1 minute,
  - B can in 2 minutes,
  - C can in 4 and
  - D can in 5 minutes.
- Unfortunately, not more than two persons can go through the narrow tunnel at one time, moving at the speed of the slower one.
- They have a single torch to show their way in dark
- What is the minimum time for all to make it to the other side of the tunnel?