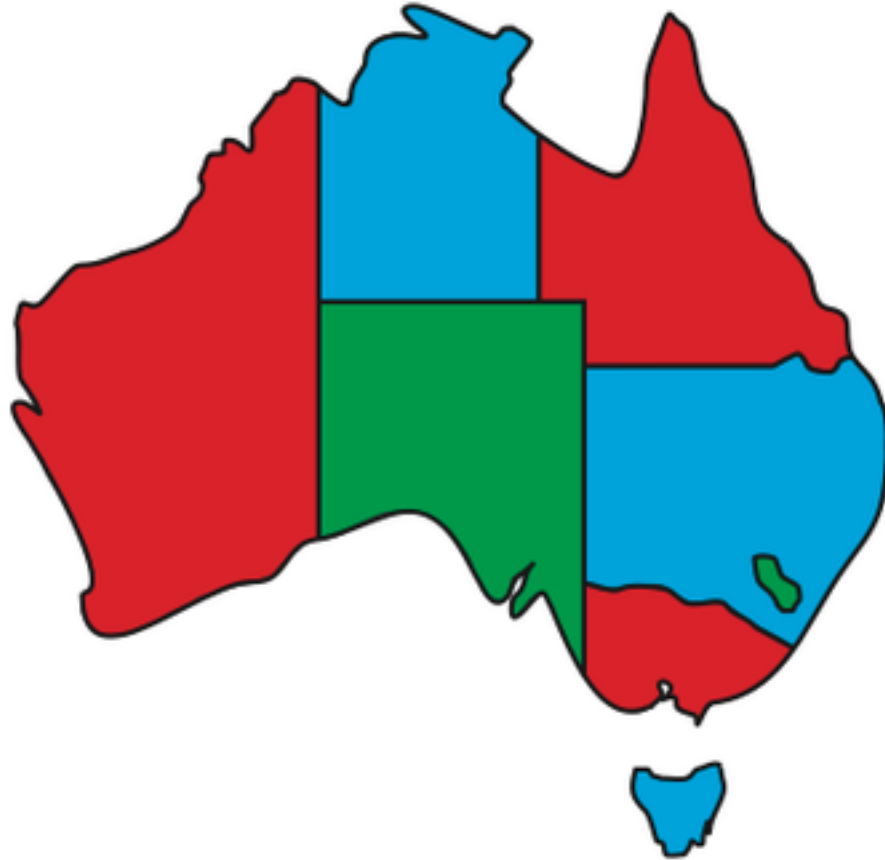
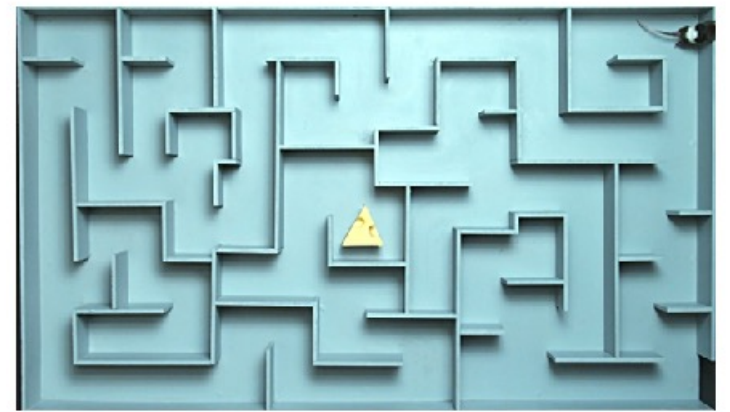
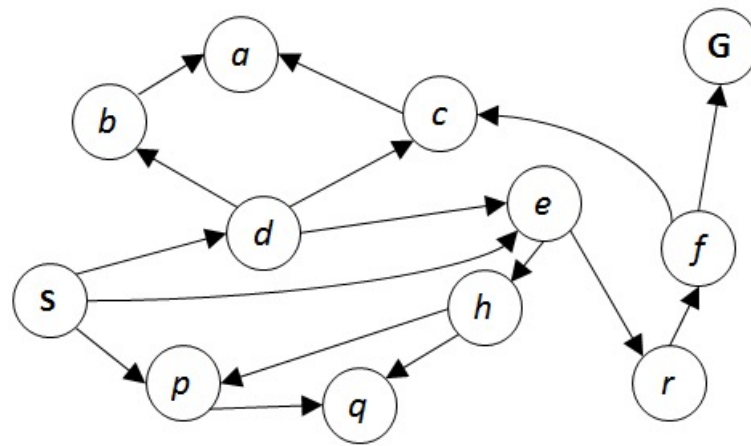


COMP 341 Intro to AI



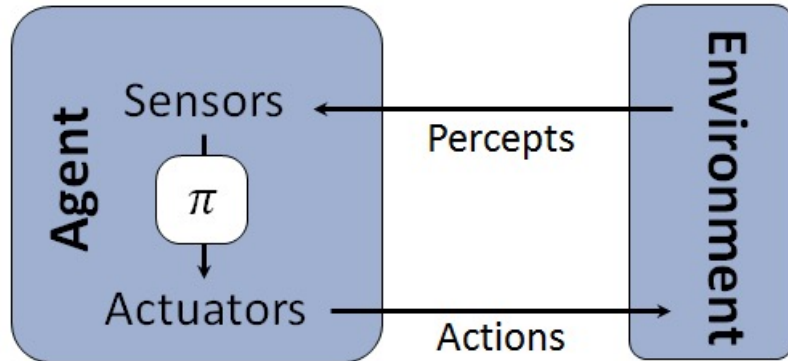
Constraint Satisfaction Problems



Asst. Prof. Barış Akgün
Koç University



Previously on Intro to AI

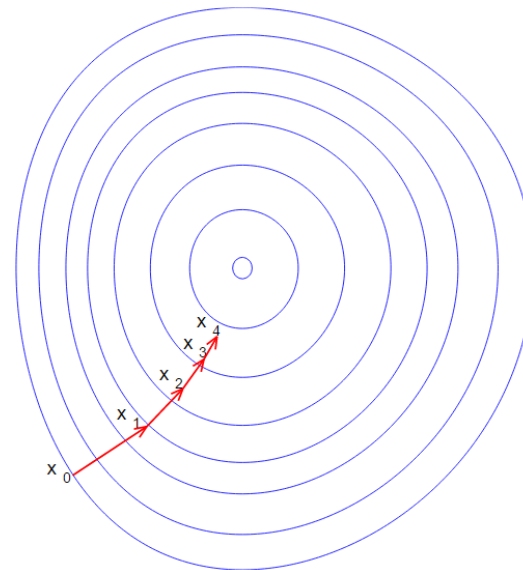
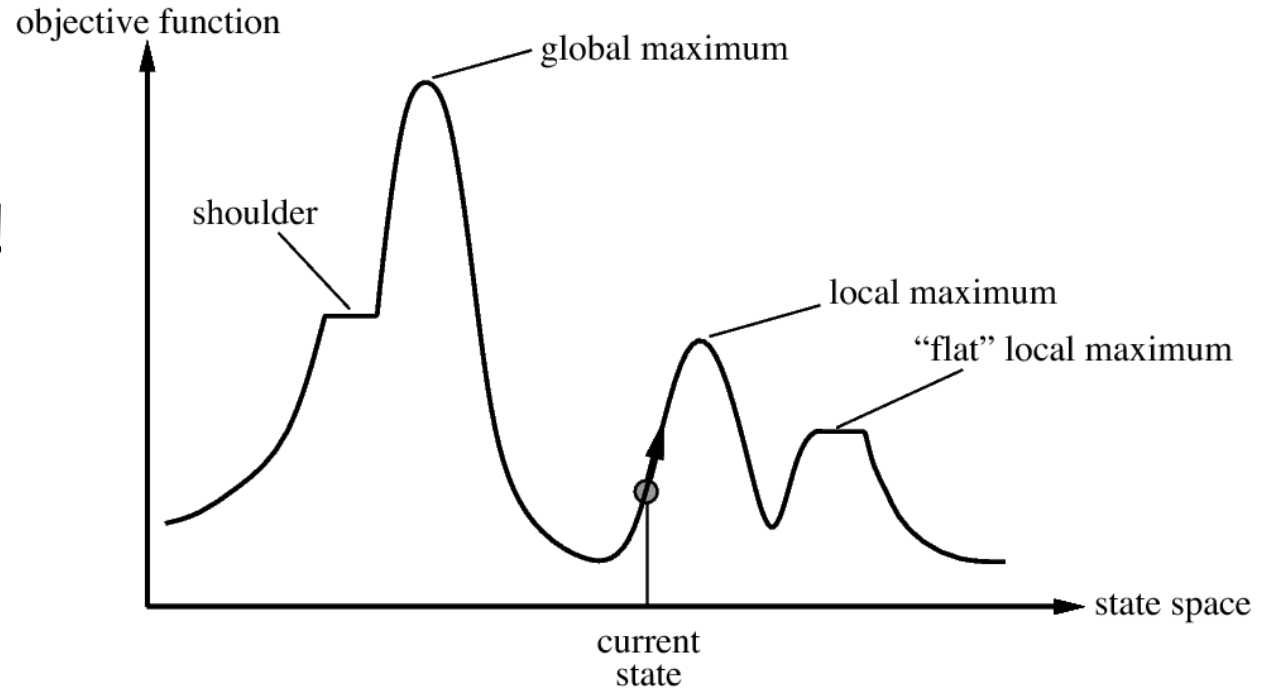


Search

- Uninformed
- Informed
- Solution is a path to goal

Local Search

- Solution is important, not the path!
- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithms
- Gradient Descent

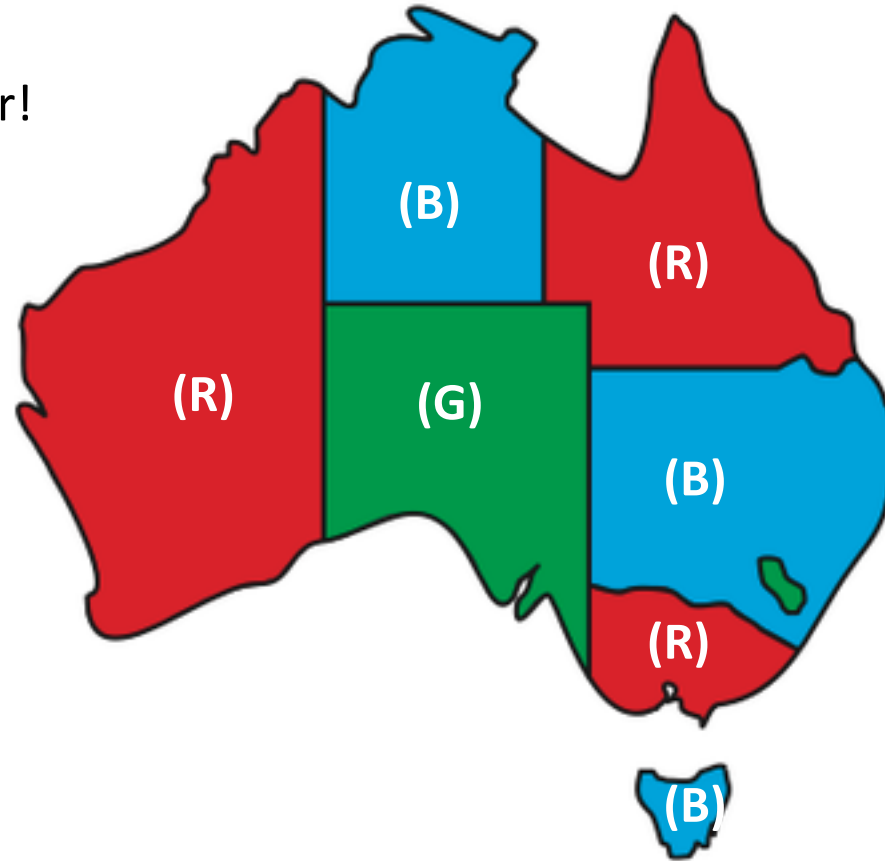


Local Search

- Formulation:
 - Current State
 - Transition Function
 - Evaluation Function and State Space “Landscape”
- Algorithms: Move towards Better States (Where the “Local” comes from)
 - Complete: Find a solution if one exists
 - Optimal: Find the best state
- Usually easy to code!

Constraint Satisfaction Problems

No neighbors with the same color!



Search So Far


- Classical Search:
 - Solution is path to a goal state
- Local Search:
 - Solution is the goal state itself
- CSPs?
 - Goal matters
 - States and goal test have specific structure!
 - Allows for general heuristics

Constraint Satisfaction Problems

- Standard Search
 - State is a **black box** data structure
 - Goal test: Can be **any Boolean function** of states
 - Successor Function: Can be **anything** that returns valid states
 - Heuristic function: Can be **anything** that maps states to a non-negative scalar
- CSPs
 - State is defined by **variables** X_i with **values** from **domain** D_i .
 - Map Coloring Example: Variables are the color of each Australian state and the domain is the set of allowable colors
 - Goal test is **a set of constraints** specifying **allowable combinations of values** for subsets of variables
 - Map Coloring Example: All states are colored and neighboring states do not have the same color
- This structure allows useful **general-purpose** algorithms with more power than standard search algorithms

CSPs

- State is defined by **variables** X_i with **values** from **domain** D_i



Domains of variables can be different!

- Goal test is **a set of constraints** specifying **allowable combinations of values** for subsets of variables. E.g.

$$\sum_{i \in A} X_i == k \quad X_i \neq X_j \text{ for } i \neq j, i \in A, j \in A$$

Real Life Example (!) - Carpool

- Ahmet, Elif, Mehmet, Zeynep want to carpool to Bolu
 - Variables are A,E,M,Z
- There are only 2 cars
 - Domains = $\{C_1, C_2\}$
- The cars belong to Ahmet and Zeynep
 - Constraints: $A = C_1$ and $Z = C_2$
- Ahmet and Elif do not like each other
 - Constraint: $A \neq E$
- Mehmet has a crush on Zeynep
 - Constraint: $M = Z$
- A solution
 - $A=C_1$, $E=C_2$, $M= C_2$, $Z=C_2$

Side Note: They should just sit Elif in the front and Ahmet in the back and take 1 car

But the problem does not model that!

Solving CSPs

- Each state of the problem is a possible assignment to some or all the variables
- **Legal Assignment:** no violations
- **Complete Assignment:** every variable assigned

Map Coloring


- Color the map such that no two neighbors have the same color
- Variables:
 - WA, NT, Q, NSW, V, SA, T
- Domains:
 - $D_i = \{\text{red, green, blue}\}$
- Constraints: adjacent regions must have different colors
 - Implicit: $WA \neq NT$
 - Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}$



Map Coloring

- Color the map such that no two neighbors have the same color
- Solutions are assignments satisfying all constraints, e.g.,

{WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green}

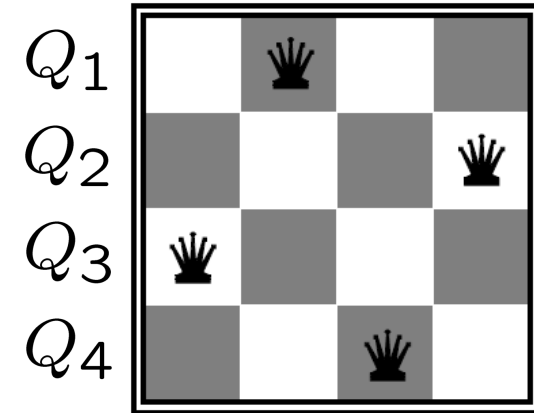


Example: N-Queens

- Variables: Q_k
- Domains: $\{1, 2, \dots, N\}$
- Constraints:

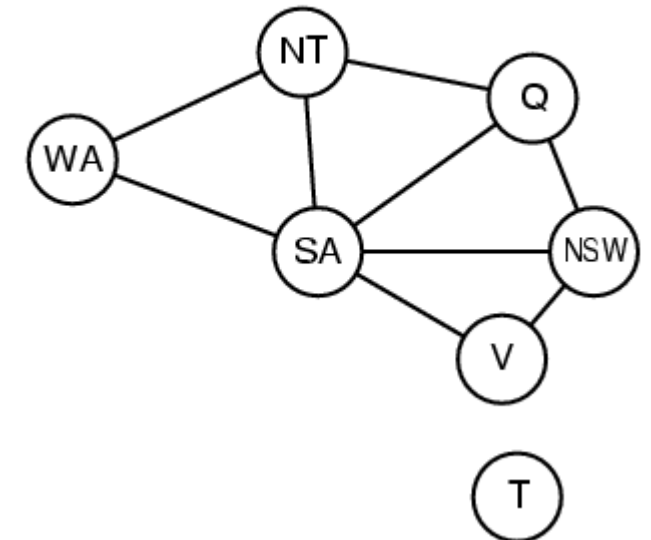
Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$
...



Constraint Graph

- Binary CSP: each constraint relates (at most) two variables
- **Binary Constraint Graph** is a data structure we use to represent the problem
 - Nodes are variables
 - Arcs show which variables are constrained



Example: Cryptarithmic

- Variables:

$F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

- Domains:

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

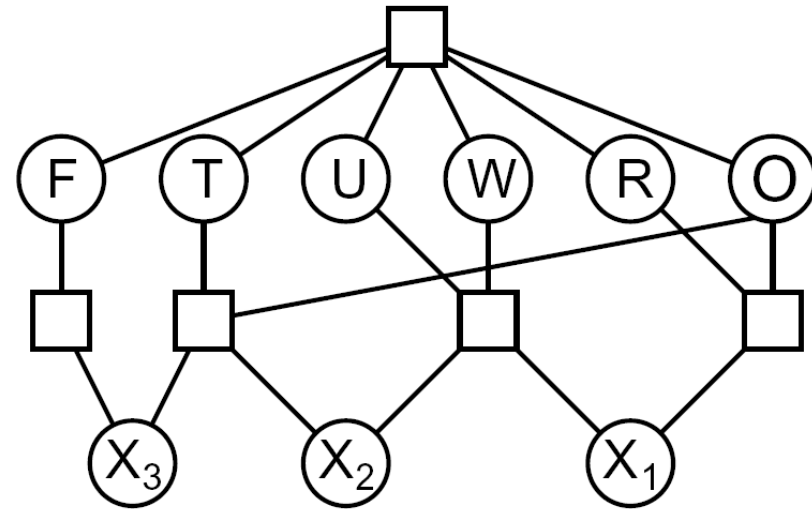
- Constraints:

$\text{alldiff}(F, T, U, W, R, O)$

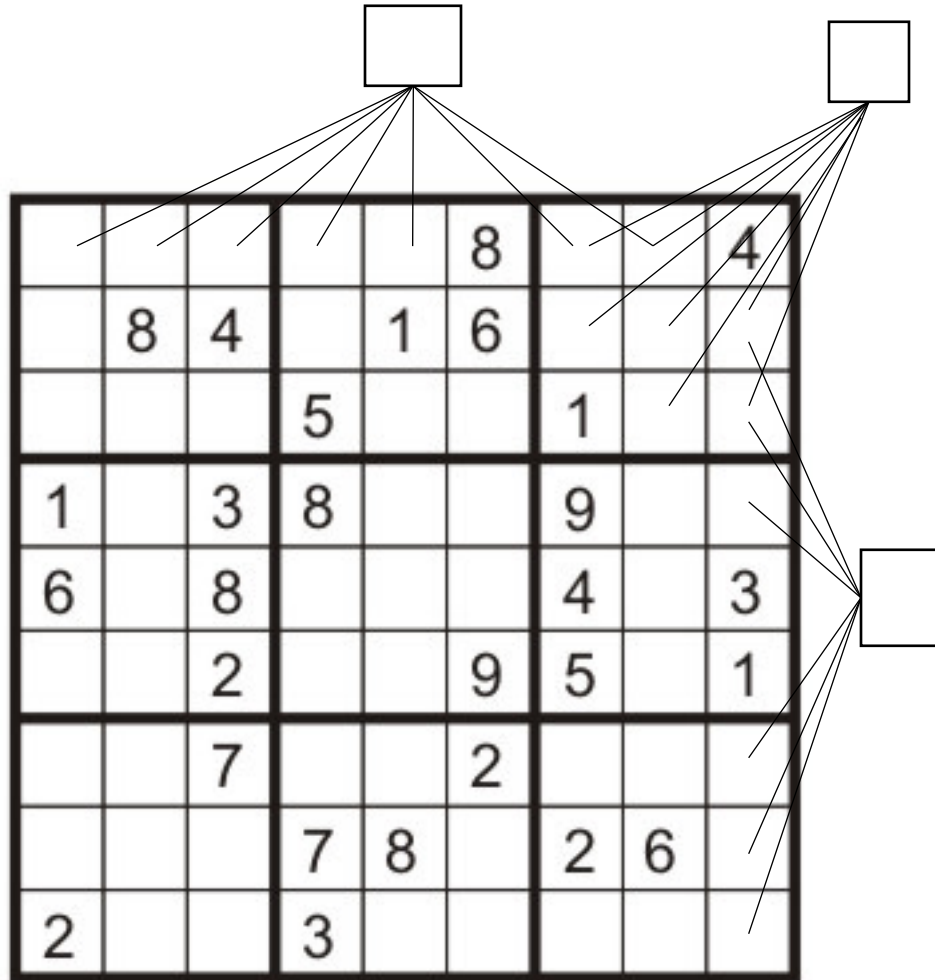
$O + O = R + 10 \cdot X_1$

\dots

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$



Most Famous CSP - Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1,2,\dots,9\}$
- Constraints:

9-way alldiff for each column

9-way alldiff for each row

9-way alldiff for each region

(or can have a bunch of
pairwise inequality
constraints)

Other Real-World Examples

- Assignment problems: e.g., who teaches what class
 - Timetabling problems: e.g., which class is offered when and where?
 - Hardware configuration
 - Transportation scheduling
 - Factory scheduling
 - Floor Planning
 - ...
-
- Many real-world problems involve real-valued variables...

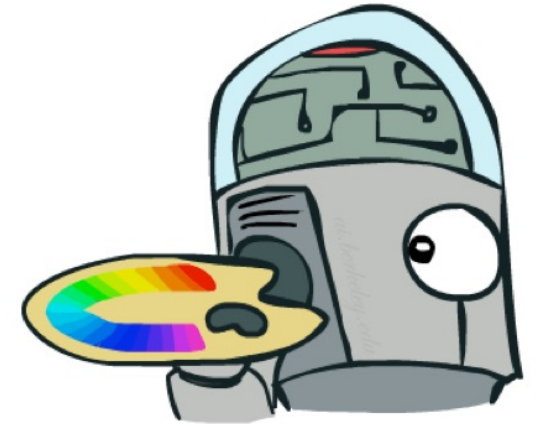
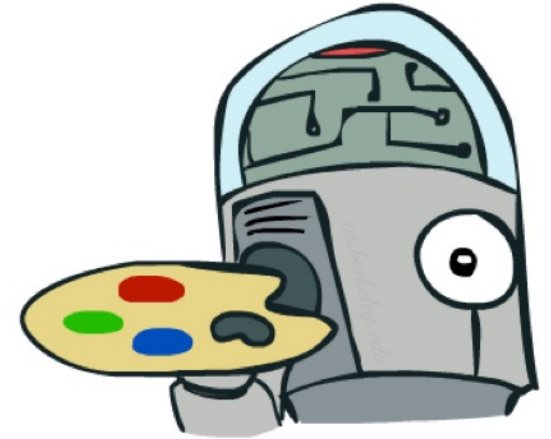
Varieties of CSPs

- Discrete Variables

- Finite domains
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
- Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Need constraint language: $Job1 + 5 < Job2$
 - Linear constraints solvable, nonlinear undecidable

- Continuous variables

- E.g., start/end times for Hubble Telescope observations
- Linear constraints solvable in polynomial time by Linear Programming methods (ever heard of the Simplex Method?)

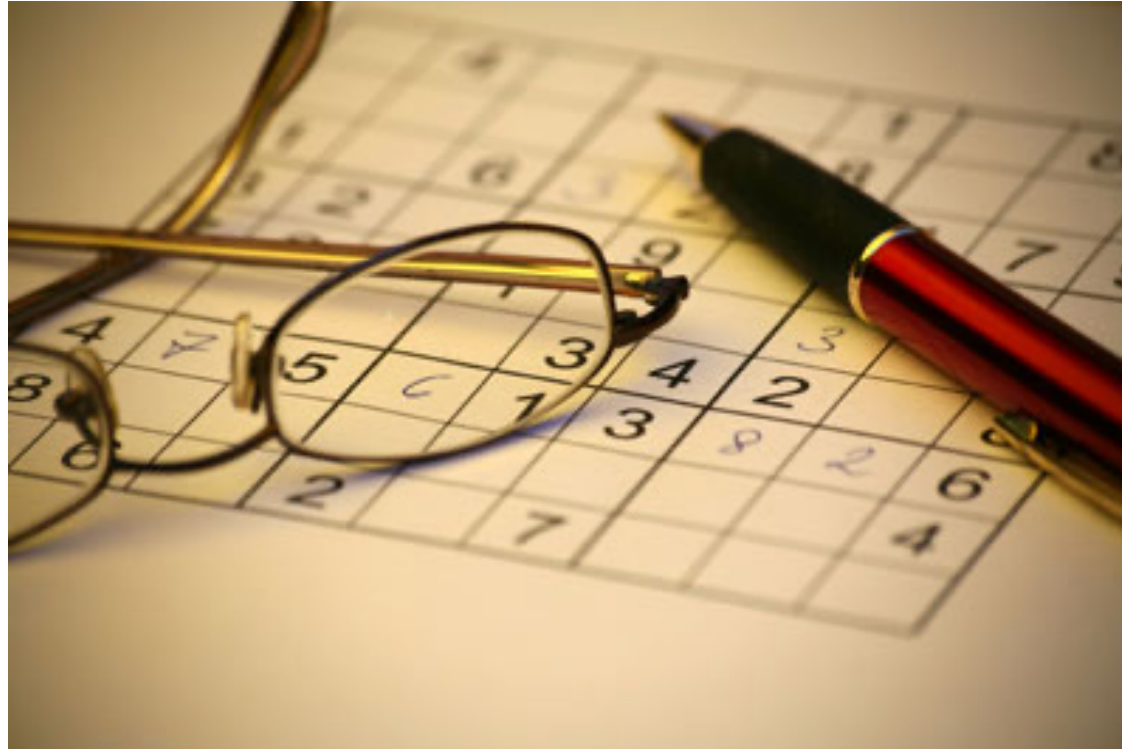


Varieties of Constraints

- Varieties of Constraints
 - Unary constraints involve a single variable (equivalent to reducing domains)
e.g.: $SA \neq green$
 - Binary constraints involve pairs of variables,
e.g.: $SA \neq WA$
 - Higher-order constraints involve 3 or more variables:
e.g.: cryptarithmic column constraints
- Preferences (soft constraints):
 - E.g., red is better than green
 - Often representable by a cost for each variable assignment
 - Gives rise to constrained optimization problems
 - More involved methods out of our scope but local search can be used (how?)



Solving CSPs



Search Formulation

Local Search Formulation

Search Formulation for CSPs

- Initial State: {}
- Successor(): assign a value (consistent with constraints) to an unassigned variable
- Goal Test(): All variables are assigned and all constraints are satisfied
- Failure: No legal assignment to do
- This is the **same** for all CSPs!
- Path is irrelevant
- Every solution appears at depth n with n variables
 - DFS anyone
- Complexity (n vars, d values)
 - Branch factor: $(n-l)d$ at depth l
 - $n!d^n$ leaves!

Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are **commutative**, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which **do not conflict** previous assignments
 - Might have to do some computation to check the constraints
 - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$

Detour: Recursive DFS

function RECURSIVE-DFS(*problem*) **returns** a solution, or failure
 return RECURSIVE-DFS_(MAKE-NODE(*problem*.INITIAL-STATE), *problem*)

function RECURSIVE-DFS_(*node*,*problem*) **returns** a solution, or failure
 if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)
 for each *action* in *problem*.ACTIONS(*node*.STATE) **do**
 child ← CHILD-NODE(*problem*, *node*, *action*)
 result ← RECURSIVE-DFS_(*child*, *problem*)
 if *result* != failure **then return** *result*
 return failure

Backtracking Search

function BACKTRACKING-SEARCH(*csp*) **returns** a solution, or failure
 return BACKTRACK({ }, *csp*)

function BACKTRACK(*assignment* , *csp*) **returns** a solution, or failure
 if *assignment* is complete **then return** *assignment*
 var ← SELECT-UNASSIGNED-VARIABLE(*csp*)
 for each value **in** ORDER-DOMAIN-VALUES(*var* , *assignment* , *csp*) **do**
 if *value* is consistent with *assignment* **then**
 add {*var* = *value*} to *assignment*
 inferences ← INFERENCE(*csp*, *var* , *value*)
 if *inferences* ≠ failure **then**
 add *inferences* to *assignment*
 result ← BACKTRACK(*assignment* , *csp*)
 if *result* ≠ failure **then**
 return *result*
 remove {*var* = *value*} and *inferences* from *assignment*
 return failure

select a
variable

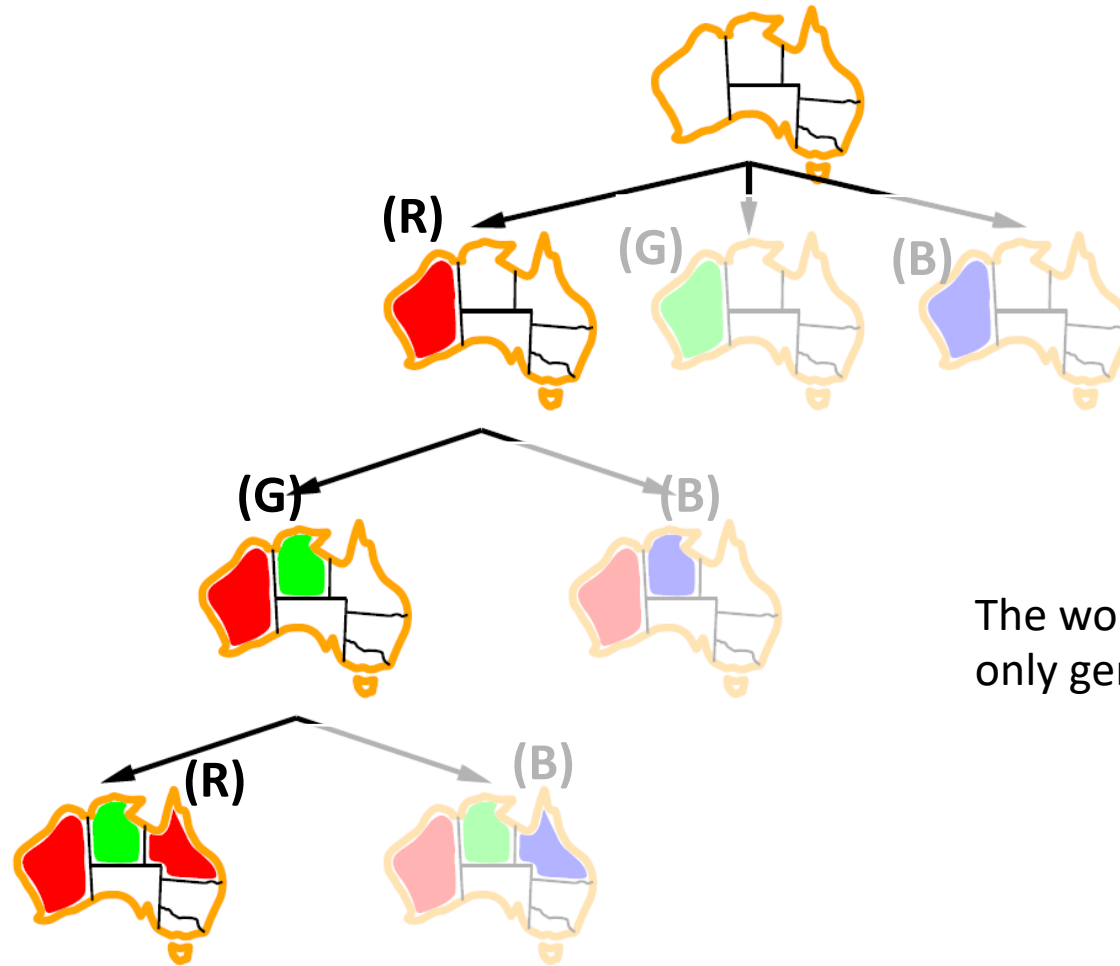
Find a value
consistent with
constraints

Recurse to
assign another

only keeps a single
representation of
the assigned state!

If no consistent assignment exists, return
failure, which causes another value to be
tried


Backtracking Example



The worn-out states are only generated if needed

Another Backtracking Example

- Variable Order: Left to right, top to bottom
- Value Order: Random (assume we have the following)
 - Red
 - Red
 - Green
 - Blue
 - Green
 - Red
 - Blue
 - Green
 - Green
 - Red
 - Blue



Improvements

- Backtracking: DFS + variable ordering + constraint checking
- Uninformed: Add heuristics to improve
- **General Purpose Heuristics** thanks to the structure of CSPs
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Filtering: What inference can be made to detect failures early ?
- Structure: Can we exploit the problem structure?

Emphasis Slide

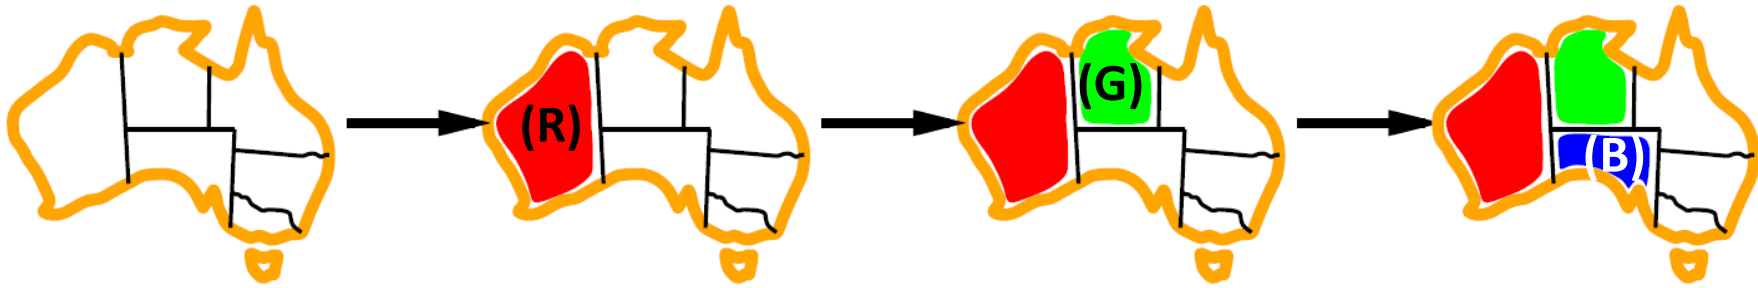
- Variables (e.g. map location) can take values (e.g. specific color) from their domain (e.g. set of colors)
- Ordering
 - Picking the **variable** to assign next
 - Picking the **value** to assign to the chose variable
- Filtering: Filter the domains
 - Removing the infeasible values from the domains(e.g. remove colors)

Variable Ordering: What variable to assign next?

- Fixed order
- Random
- Other ideas?
 - Let's look at the number of constraints per variable!

Minimum Remaining Values (MRV)


- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal left values in its domain



- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering + not running out of options

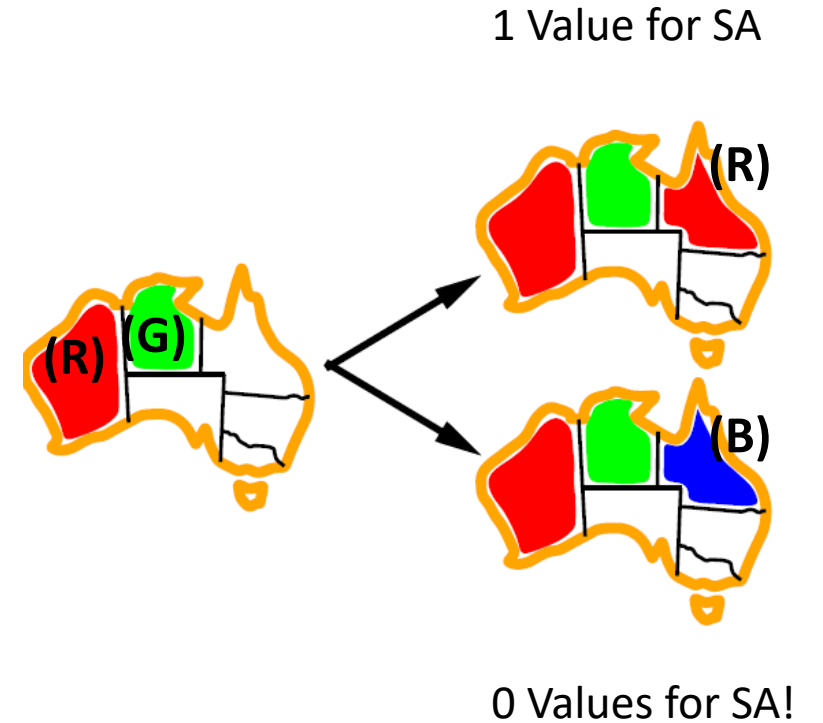
Degree Heuristic

- Tie breaker among MRV variables
- Choose the variable with most constraints on remaining variables



Least Constraining Value

- Which value to assign next?
- Value Ordering: Least Constraining Value
 - Given a choice of variable, choose the **least constraining value**
 - I.e., the one that **rules out the fewest values** in the remaining variables
 - Note that it may take some computation
- Why least rather than most?
- Combining these ordering ideas makes 1000 queens feasible




Summary of Ordering

- Which variable to select next and which value to assign to it?
- Detect failures early (MRV + DH) – Picking variables
- Enter the most promising branch (LCV) – Picking values
- Note that the heuristics do not change the theoretical bounds!

Emphasis Slide

- Variables (e.g. map location) can take values (e.g. specific color) from their domain (e.g. set of colors)
- Ordering
 - Picking the **variable** to assign next – MRV (+ DH as tie breaker)
 - Picking the **value** to assign to the chose variable - LCV

Backtracking + Heuristics



- Variable Domains
 - WA: {**R**, **G**, **B**}
 - NT: {**R**, **G**, **B**}
 - SA: {**R**, **G**, **B**}
 - Q: {**R**, **G**, **B**}
 - NSW: {**R**, **G**, **B**}
 - V: {**R**, **G**, **B**}
 - T: {**R**, **G**, **B**}
- Tie Breaking Variable Order up-down, left-right from above
- Tie Breaking Value Order: **R** – **G** – **B**



Step	MRV	DH	LCV	Assignment
1				
2				
3				
4				
5				
6				
7				

Backtracking + Heuristics



- Variable Domains

- WA: {~~R~~, ~~G~~, ~~B~~}
- NT: {~~R~~, ~~G~~, ~~B~~}
- SA: {~~R~~, ~~G~~, ~~B~~}
- Q: {~~R~~, ~~G~~, ~~B~~}
- NSW: {~~R~~, ~~G~~, ~~B~~}
- V: {~~R~~, ~~G~~, ~~B~~}
- T: {~~R~~, ~~G~~, ~~B~~}

- Tie Breaking Variable Order up-down, left-right from above

- Tie Breaking Value Order: ~~R~~ – ~~G~~ – ~~B~~

Step	MRV	DH	LCV	Assignment
1	Same (3)	SA (5)	Same	SA - R
2	All (2) but T (3)	NT, Q, NSW (2)	Same	NT - G
3	WA, Q (1)	Q (1)	Only B	Q - B
4	WA, NSW (1)	NSW (1)	Only G	NSW - G
5	WA, V (1)	Same (1)	Only B	WA - B
6	V (1)	-	Only B	V - B
7	T (3)	-	Same	T - R



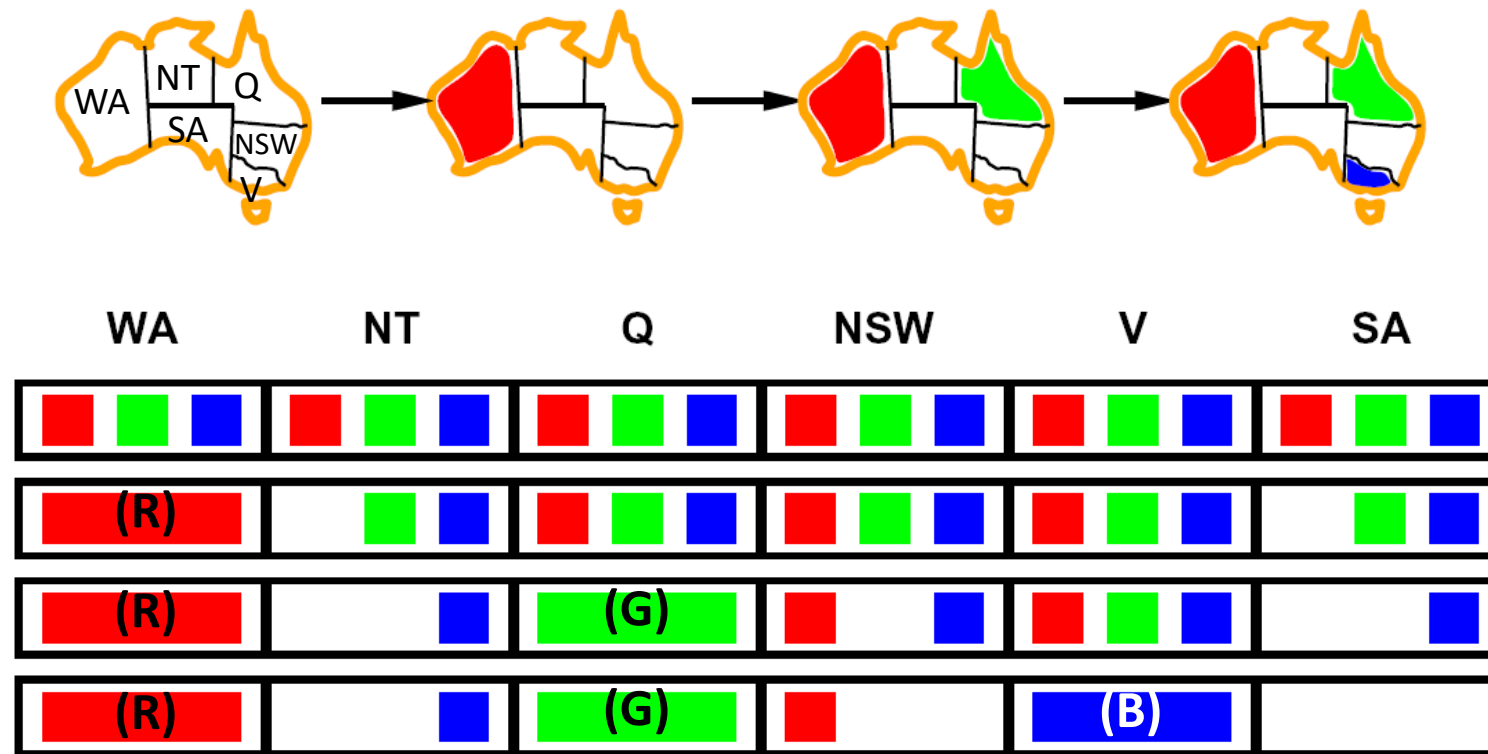
We didn't need to backtrack for this example, but this is not the norm

Filtering

- How to detect failures early?
- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward Checking
- Constraint Propagation - Arc consistency

Forward Checking


- Forward checking: Cross off values that violate a constraint when added to the existing assignment
- Backtrack when no assignments left



MRV + Forward Checking: FC can be used to compute what MRV needs!

Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures
- After deleting neighbors, check constraints for all other variables





WA	NT	Q	NSW	V	SA
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div>(R)</div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div>(R)</div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

- NT and SA cannot both be blue!
- Why didn't we detect this yet?
- *Constraint propagation*: reason from constraint to constraint

Legend:
Left: Red
Middle: Green
Right: Blue

Consistency of A Single Arc

- An arc $X \rightarrow Y$ is **consistent** iff for *every* x in the tail there is *some* y in the head which could be assigned without violating a constraint




Legend:
Left: Red
Middle: Green
Right: Blue

- Delete from tail!
- Forward checking: Enforcing consistency of arcs pointing to each new assignment

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


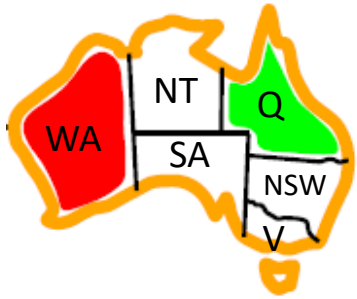


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


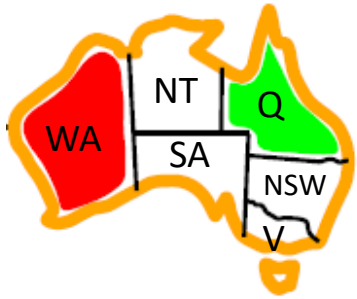


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:





Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


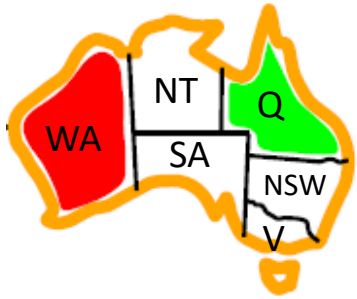


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


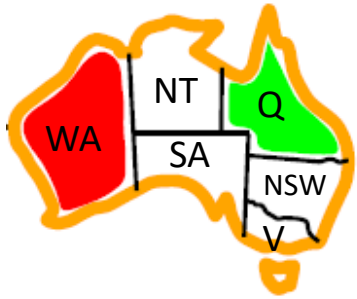


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


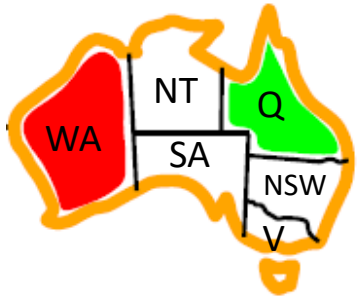


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


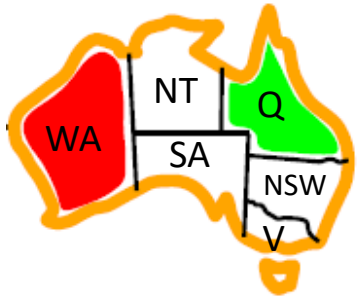


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


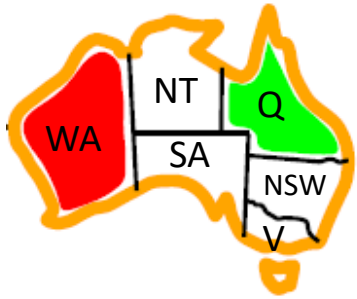


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:


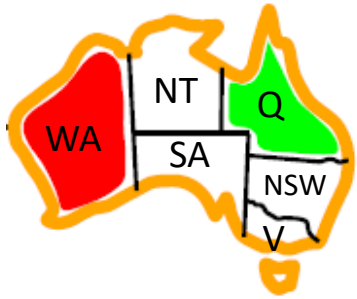


Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking
- Can be run as before or after each assignment
- What's the downside of enforcing arc consistency?

Legend:
Left: Red
Middle: Green
Right: Blue

*Remember: Delete
from the tail!*

Enforcing Arc Consistency in a CSP

Check consistency and
remove value if necessary

Add all the neighbors to
the queue if something is
removed

Delete from tail to
enforce consistency

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X, D, C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REVISE(*csp*, X_i, X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k in $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

add (X_k, X_i) to queue

return true

function REVISE(*csp*, X_i, X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow false

for each x in D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

delete x from D_i

revised \leftarrow true

return revised

- Runtime: $O(n^2d^3)$, can be reduced to $O(n^2d^2)$
- ... but detecting all possible future problems is NP-hard

Arc Consistency and Forward Checking


- FC is essentially an arc consistency check for a single node!
 - Head is the assigned node, tails are its neighbors
- If you ran AC, no need to run FC
- FC is faster per value-assignment
- AC can catch failures earlier

Summary of Filtering

- FC: Remove values from the domains neighboring nodes
 - Fast to compute
 - Plays well with MRV
 - Does not catch some failures early
- Arc Consistency: Make all arcs consistent after an assignment
 - Keep checking arcs until there is no change
 - Entails FC
 - Earlier failure detection
 - Costly to Compute (especially if there are a lot of constraints)


Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



It's rare to find unconnected parts of the graph


Tree-Structured CSPs



- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time
- Compare to general CSPs, where worst-case time is $O(d^n)$

(Skipping the algorithm this semester)

Nearly Tree-Structured CSPs



- **Conditioning:** instantiate a variable, prune its neighbors' domains
- **Cutset conditioning:** instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size c gives runtime $O((d^c) (n-c) d^2)$, very fast for small c


Cutset Conditioning

Choose a cutset

Instantiate the cutset
(all possible ways)

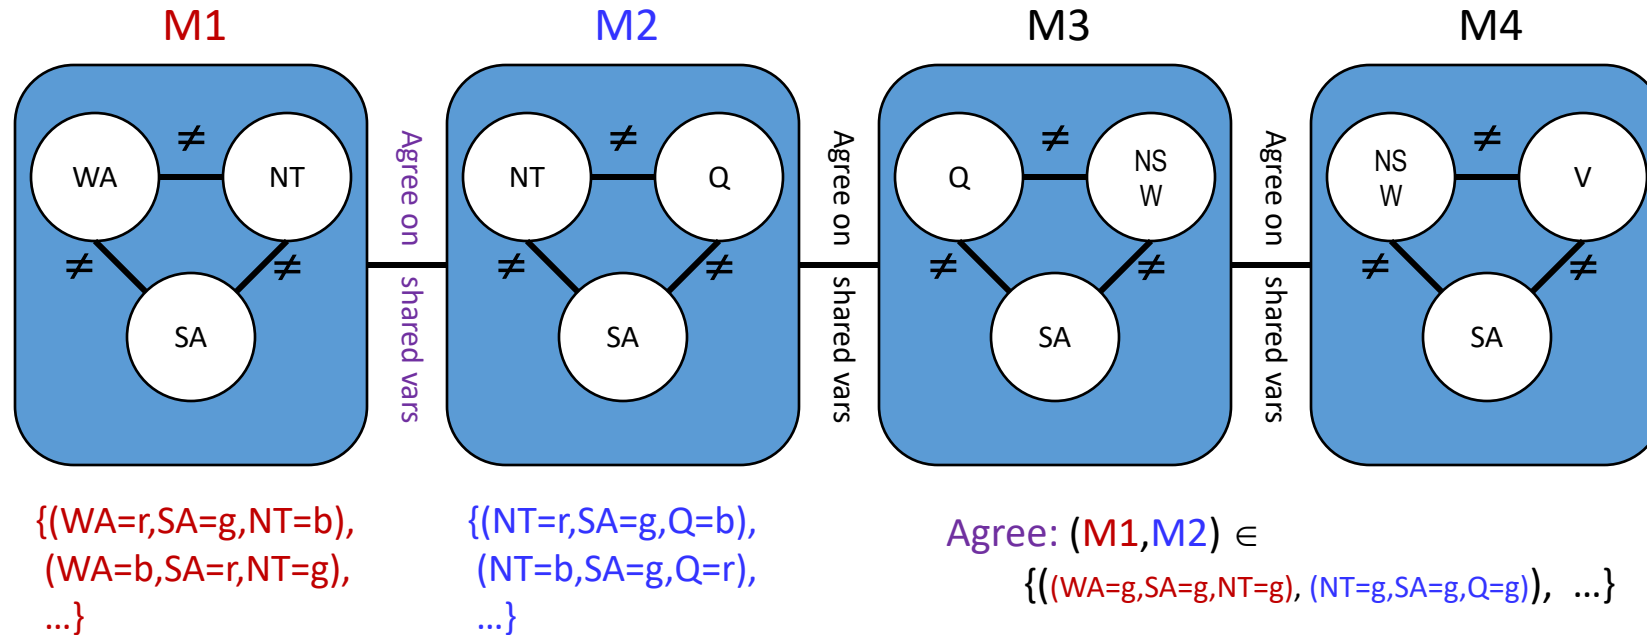

Compute residual CSP
for each assignment

Solve the residual CSPs
(tree structured)



Tree Decomposition*

- Idea: create a tree-structured graph of mega-variables
- Each mega-variable encodes part of the original CSP
- Subproblems overlap to ensure consistent solutions



Local Search For CSPs – MIN-CONFLICTS

function MIN-CONFLICTS(*csp*, *max_steps*) **returns** a solution or failure

inputs: *csp*, a constraint satisfaction problem

max_steps, the number of steps allowed before giving up

current \leftarrow an **initial complete assignment** for *csp*

for $i = 1$ to *max_steps* **do**

if *current* is a solution for *csp* **then return** *current*


$var \leftarrow$ a **randomly chosen conflicted variable** from *csp*.VARIABLES

$value \leftarrow$ the value v for *var* that **minimizes CONFLICTS**(*var*, v , *current*, *csp*)

 set *var* = *value* in *current*

return failure

Example: 4-Queens





- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n)$ = number of attacks

Performance of Min-Conflicts

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



Related to the “Phase Transition” phenomenon

Summary of CSPs

- CSPs are a special kind of search problem:
 - States are partial assignments
 - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
 - Ordering
 - Filtering
 - Structure
- Iterative min-conflicts is often effective in practice